

# Exploratory Data Analysis and Topic Modeling

This section will cover:

- initial exploratory analysis
- grammar and spelling correction
- feature extraction with SpaCy
- topic modeling with latent Dirichlet allocation (LDA).
- automatic scoring with LDA

```
In [1]: #essential imports  
%matplotlib inline  
  
import numpy as np  
import pandas as pd  
import re  
from datetime import datetime
```

```
In [2]: #visualization  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [3]: #Text Processing Libraries  
import spacy  
from spacy.lang.en.stop_words import STOP_WORDS  
from string import punctuation
```

```
In [4]: #ML Libraries  
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.decomposition import LatentDirichletAllocation  
from sklearn.model_selection import train_test_split
```

```
import joblib
```


```
In [5]: #LDA Visualization
import pyLDAvis.lda_model
from pyLDAvis.lda_model import prepare
```

```
In [6]: #matplotlib and pandas config
import seaborn as sns
#plt.style.use('seaborn-colorblind')
palette = sns.color_palette("colorblind")
plt.gca().set_prop_cycle('color', palette)

# Setup Pandas
pd.set_option('display.width', 500)
pd.set_option('display.max_columns', 100)
pd.set_option('display.notebook_repr_html', True)
pd.set_option('display.max_colwidth', 100)

pyLDAvis.enable_notebook()
# plt.rcParams['figure.figsize'] = [8, 5]
plt.rcParams['figure.dpi'] = 100

import warnings
warnings.simplefilter("ignore", DeprecationWarning)
```



```
In [7]: training_set = pd.read_csv('training_set_rel3.tsv', sep='\t', encoding = "ISO-8859-1")\
        .rename(columns={'essay_set': 'topic', 'domain1_score': 'target_score', 'domain2_score': 'topic'})\
        training_set.sample()
```

Out [7]:

	essay_id	topic	essay	rater1_domain1	rater2_domain1	rater3_domain1	target_score	rater1_domain2	rater
	10409	16359	6	The Mooring Mast has two obstacles @CAPS1 the winds on top of the building are constantly shifti...	2	2	NaN	2	NaN



In [8]:

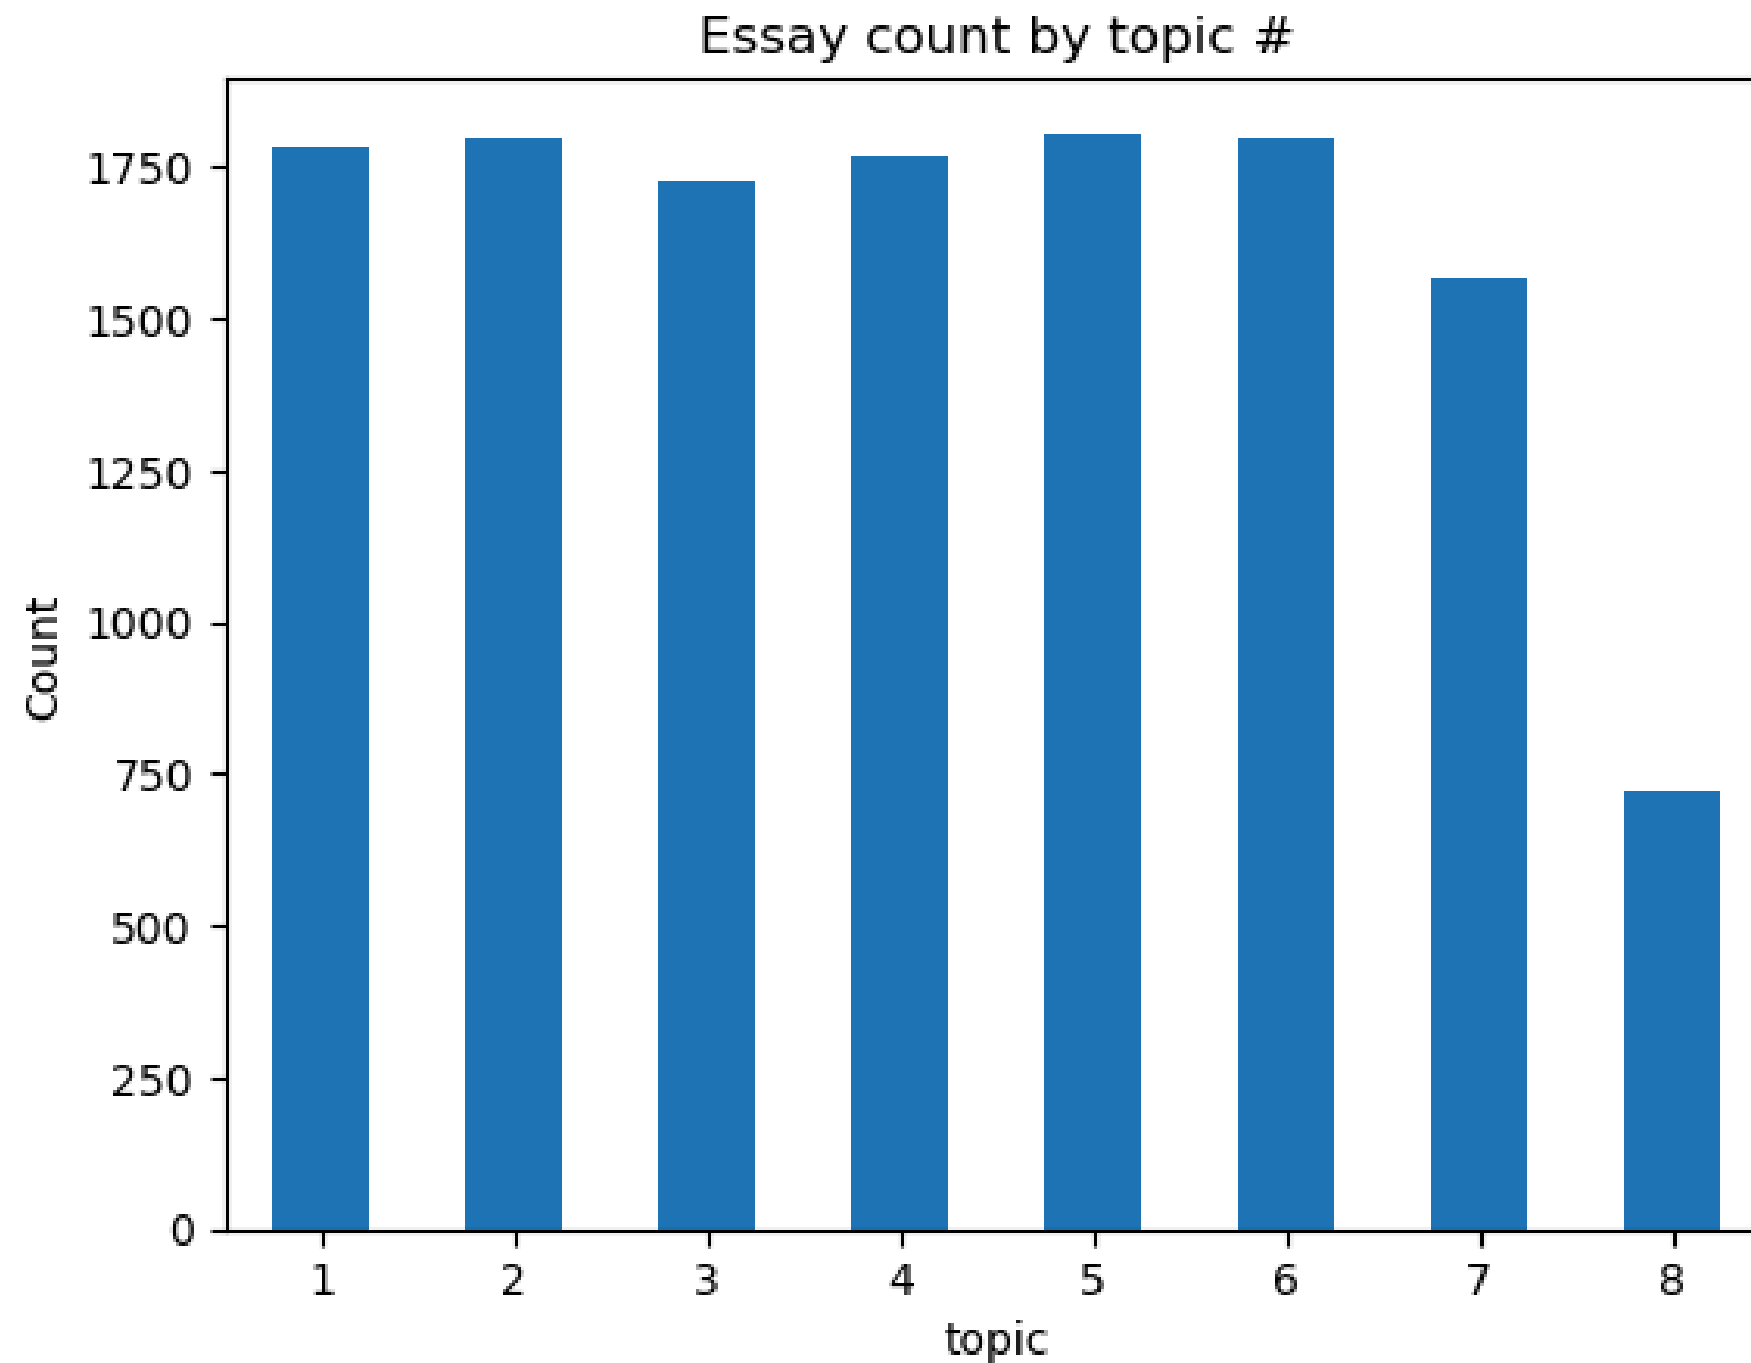
```
training_set.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12976 entries, 0 to 12975
Data columns (total 28 columns):
#   Column                Non-Null Count  Dtype
---  -
0   essay_id              12976 non-null  int64
1   topic                 12976 non-null  int64
2   essay                 12976 non-null  object
3   rater1_domain1        12976 non-null  int64
4   rater2_domain1        12976 non-null  int64
5   rater3_domain1        128 non-null    float64
6   target_score          12976 non-null  int64
7   rater1_domain2        1800 non-null   float64
8   rater2_domain2        1800 non-null   float64
9   topic2_target         1800 non-null   float64
10  rater1_trait1          2292 non-null   float64
11  rater1_trait2          2292 non-null   float64
12  rater1_trait3          2292 non-null   float64
13  rater1_trait4          2292 non-null   float64
14  rater1_trait5          723 non-null    float64
15  rater1_trait6          723 non-null    float64
16  rater2_trait1          2292 non-null   float64
17  rater2_trait2          2292 non-null   float64
18  rater2_trait3          2292 non-null   float64
19  rater2_trait4          2292 non-null   float64
20  rater2_trait5          723 non-null    float64
21  rater2_trait6          723 non-null    float64
22  rater3_trait1          128 non-null    float64
23  rater3_trait2          128 non-null    float64
24  rater3_trait3          128 non-null    float64
25  rater3_trait4          128 non-null    float64
26  rater3_trait5          128 non-null    float64
27  rater3_trait6          128 non-null    float64
dtypes: float64(22), int64(5), object(1)
memory usage: 2.8+ MB
```

```
In [9]: training_set.groupby('topic').agg('count').plot.bar(y='essay', rot=0, legend=False)
plt.title('Essay count by topic #')
```

```
plt.ylabel('Count')  
plt.show()
```




Since topic 8 has the fewest essays and the most distinct scores, it might prove to be the most challenging topic to model.

```
In [10]: # Count characters and words for each essay  
training_set['word_count'] = training_set['essay'].str.strip().str.split().str.len()
```

```
In [11]: training_set.hist(column='word_count', by='topic', bins=25, sharey=True, sharex=True, layout=(2, 4), figsize=(10, 10))  
plt.suptitle('Word count by topic #')
```

```
plt.xlabel('Number of words')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```



```
In [12]: training_set.groupby(['topic'])['target_score'].agg(['min', 'max', 'count', 'nunique'])
```

Out [12]:



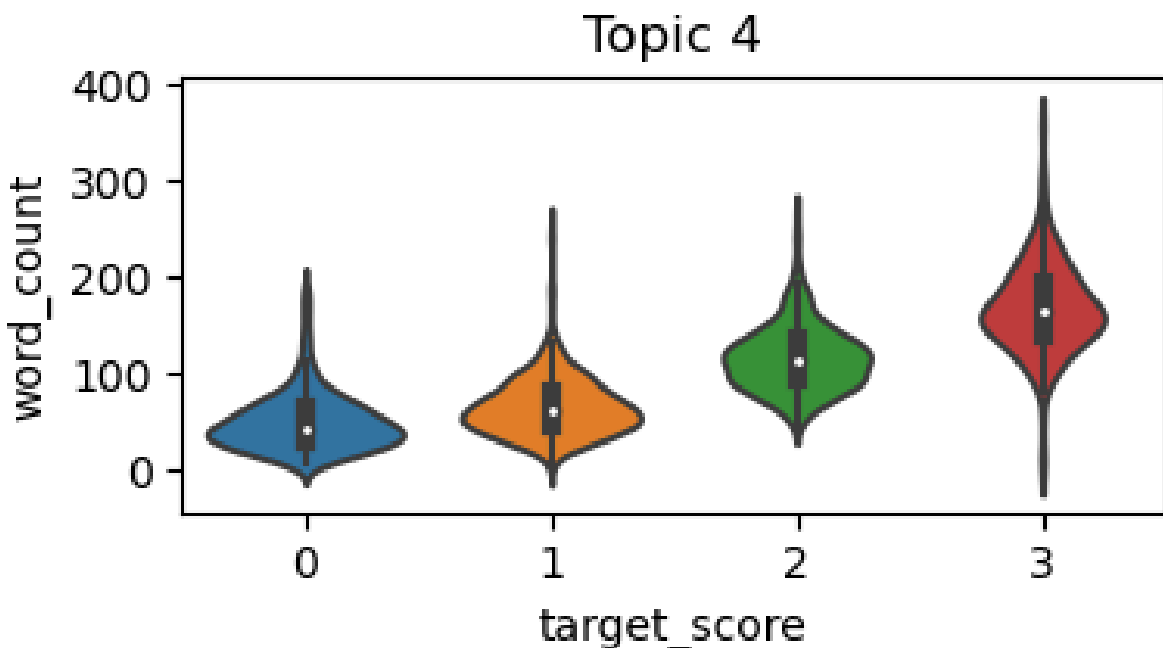
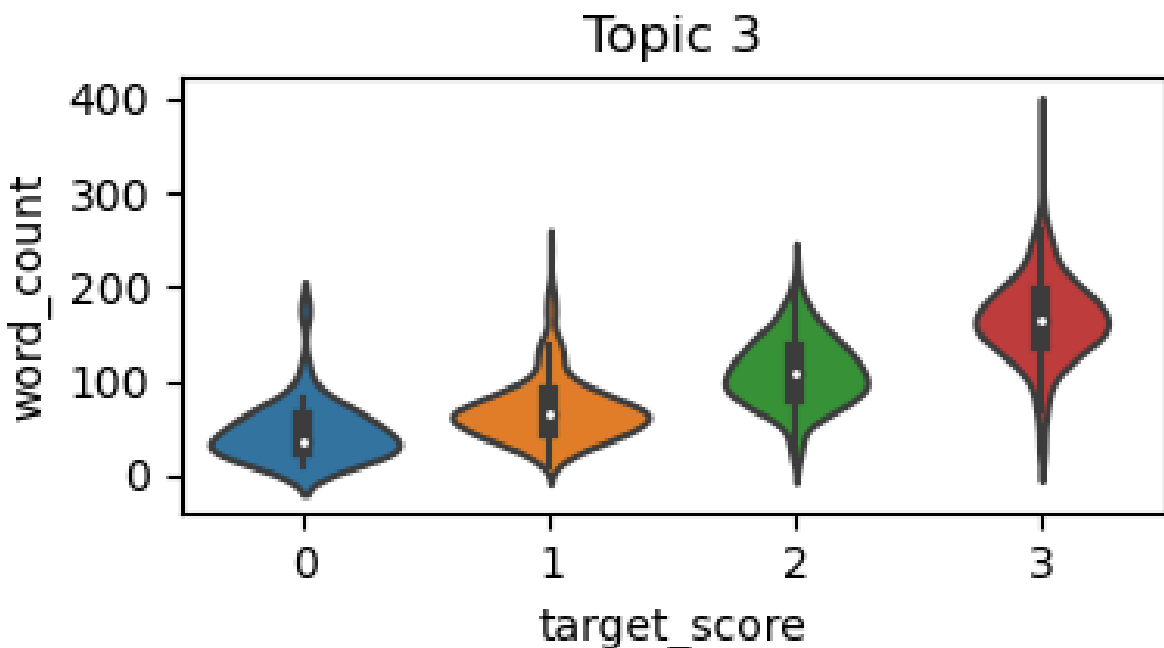


	min	max	count	nunique
--	-----	-----	-------	---------

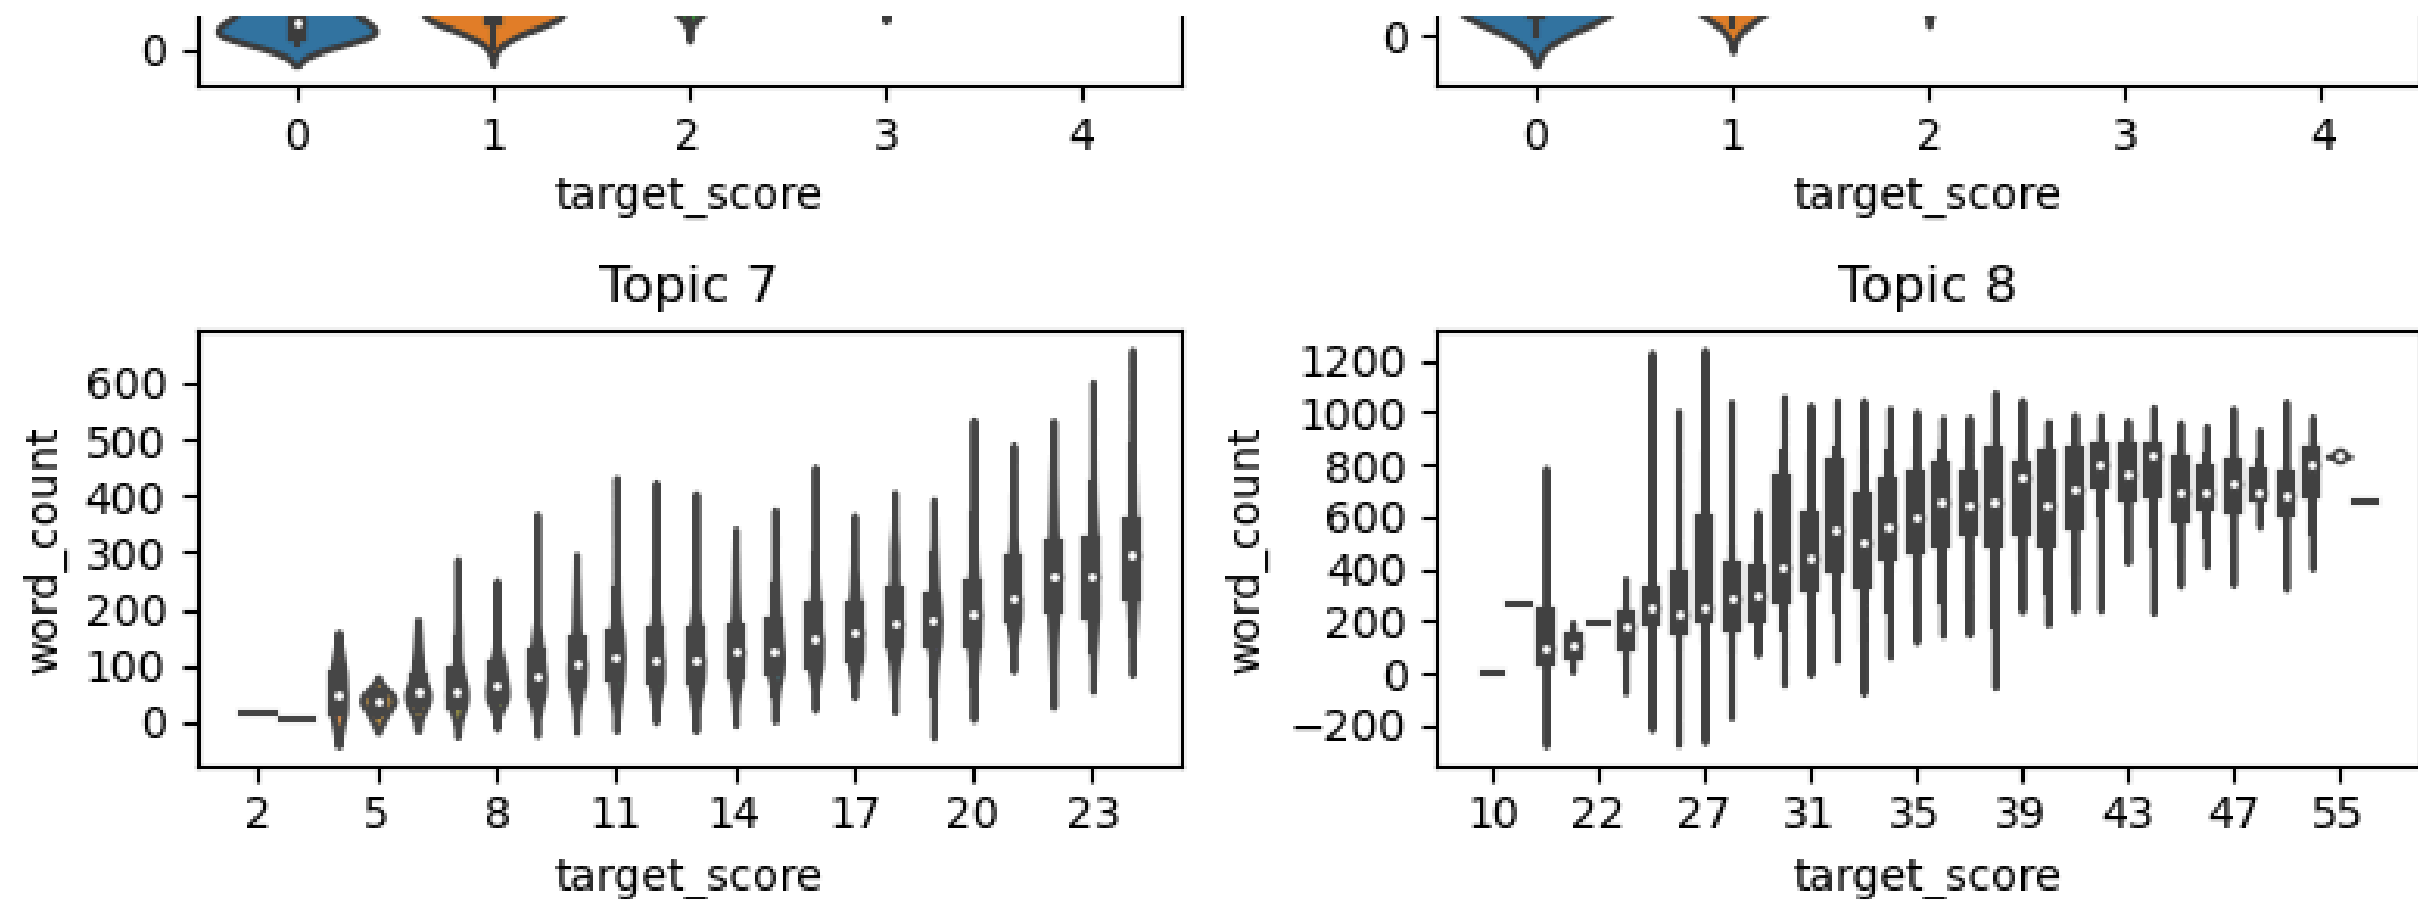
topic				
1	2	12	1783	11
2	1	6	1800	6
3	0	3	1726	4
4	0	3	1770	4
5	0	4	1805	5
6	0	4	1800	5
7	2	24	1569	23
8	10	60	723	34

```
In [13]: topic_number = 0
fig, ax = plt.subplots(4,2, figsize=(8,10))
for i in range(4):
    for j in range(2):
        topic_number += 1
        sns.violinplot(x='target_score', y='word_count', data=training_set[training_set['topic'] == topic_number])
        ax[i,j].set_title('Topic %i' % topic_number)
ax[3,0].locator_params(nbins=10)
ax[3,1].locator_params(nbins=10)
plt.suptitle('Word count by score')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```



Word count by score


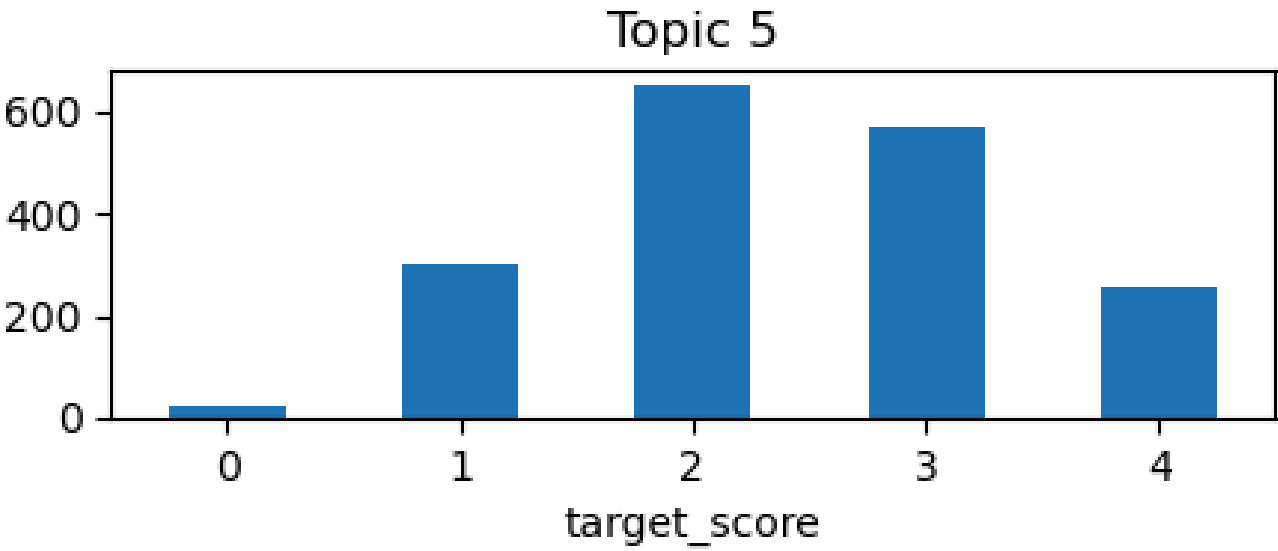
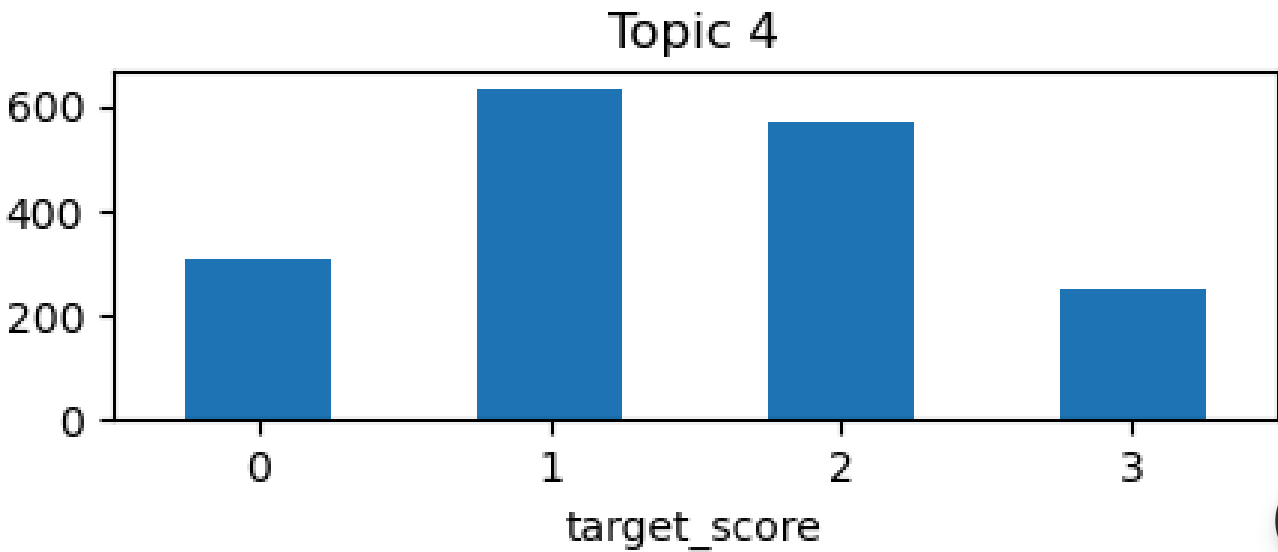
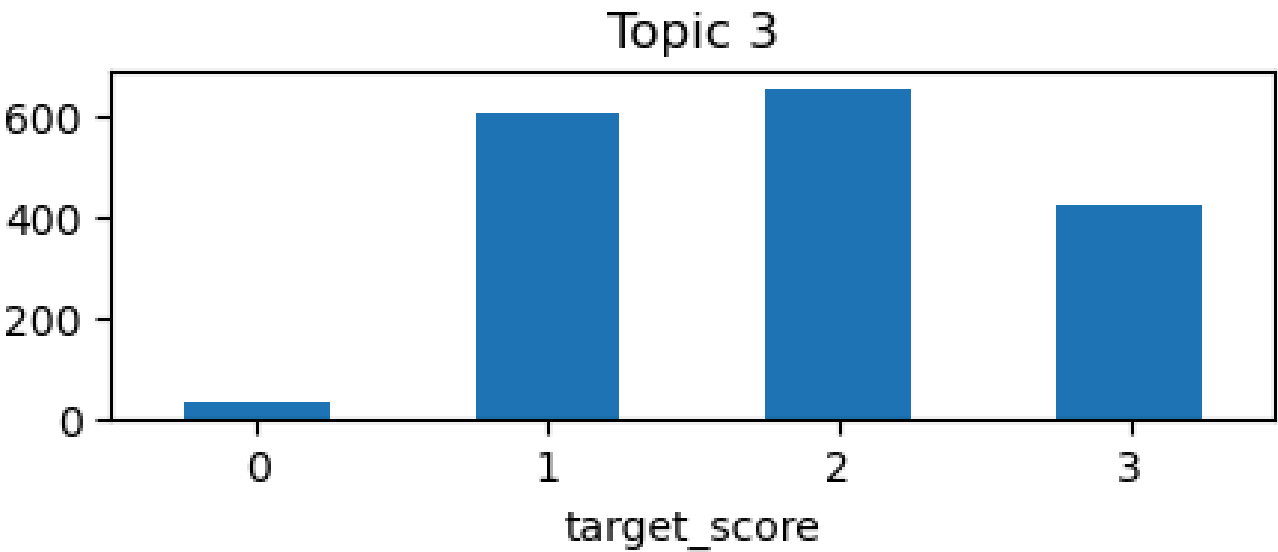







We can see a reasonable correlation between word count and score for all but topic 8 where the word count apparently reaches a maximum at the upper third of the scores.

```
In [14]: topic_number = 0
fig, ax = plt.subplots(4,2, figsize=(9,9), sharey=False)
for i in range(4):
    for j in range(2):
        topic_number += 1
        training_set[training_set['topic'] == topic_number]\
            .groupby('target_score')['essay_id']\
            .agg('count')\
            .plot.bar(ax=ax[i, j], rot=0)
        ax[i,j].set_title('Topic %i' % topic_number)
ax[3,0].locator_params(nbins=10)
ax[3,1].locator_params(nbins=10)
plt.suptitle('Histograms of essay scores')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

Histograms of essay scores





Many scores are underrepresented. Classification could be difficult without rebalancing.

## Essay processing:

1. Language correction with languagetool (add number of corrections as feature)
2. Sentence tokenization with Spacy
3. Sentence count and length
4. Word tokenize with Spacy
5. Word token classification (punctuation, stop words and anonymized entities, pos, ent)
6. Document similarity (based on arbitrary high scoring essay for each topic)

## Grammar and spelling

As to be expected with student essays, many essays exhibit grammar and spelling errors.

Parts-of-speech (POS) and named-entity-recognition (NER) is hampered in part by the lack of consistent spelling and punctuation. Therefore, the essays will be corrected using languagetool and the nlp parsing will be performed with Spacy on the corrected essays.

```
In [15]: text = 'Some people are still using Myspoce instead of facdbook'
```

Textblob is an open source nlp package and something to keep an eye on. Unlike Spacy, textblob includes a statistics based spell checker. It only claims 70% accuracy, and in my testing it didn't perform well.

<https://textblob.readthedocs.io/en/dev/>

```
In [16]: from textblob import TextBlob
data = TextBlob(text)
print (data.correct())
```

Some people are still using Myspoce instead of facdbbook

Languetool.org has a python wrapper for spelling and grammatical errors at <https://pypi.org/project/language-tool-python/>

It appears to work quite well, although it requires intermediate storage of a list of errors ('matches').

```
In [17]: import language_tool_python
tool = language_tool_python.LanguageTool('en-US')

matches = tool.check(text)
tool.correct(text)
```

```
Out[17]: 'Some people are still using MySpace instead of Facebook'
```

```
In [18]: text = training_set.essay[1871]
text
```

```
Out[18]: 'I do think that there should be a censorship in not just in libraries, but everywhere.      Personlly, I
think that the way that the libraries have the books are appropite and if the parents do not want thier ch
ildren going any where that is not privy to them keep a hand lenght away      As for the parents, the pare
nts know the aera that intrest them ,therefor the parents should go there'
```

```
In [19]: matches = tool.check(text)
tool.correct(text)
```

Out[19]: 'I do think that there should be a censorship in not just in libraries, but everywhere. Personally, I think that the way that the libraries have the books are appropriate and if the parents do not want their children going anywhere that is not privy to them keep a hand length away As for the parents, the parents know the area that interest them, therefor the parents should go there'

As seen in the last sentence, languagetool did not correct `",therefor"`. Nonetheless, it should be good enough to proceed.

```
In [20]: """
use language tool to correct for most spelling and grammatical errors. Also count the applied corrections.
Using language_check python wrapper for languagetool:
https://pypi.org/project/language-tool-python/
"""

tool = language_tool_python.LanguageTool('en-US')

t0 = datetime.now()

training_set['matches'] = training_set['essay'].apply(lambda txt: tool.check(txt))
training_set['corrections'] = training_set.apply(lambda l: len(l['matches']), axis=1)
training_set['corrected'] = training_set.apply(lambda l: tool.correct(l['essay']), axis=1)

t1 = datetime.now()
print('Processing time: {}'.format(t1 - t0))

# save work
training_set.to_pickle('training_corr.pkl')
```

Processing time: 0:27:42.024892

Here's a very special example of poor writing skills (or perhaps a digitization error?). None of the spell checkers I tried were able to make much sense out of this.

```
In [21]: print('Original:')
print(training_set.essay[18])
print('Corrected with languagetool:')
print(training_set.corrected[18])
```

Original:

I aegre waf the evansmant ov tnachnolage. The evansmant ov tnachnolige is being to halp fined a kohar froi alnsas. Tnachnolage waf ont ot we wod not go to the moon. Tnachnologe evans as we maech at. The people are in tnacholege to the frchr fror the good ov live. Famas invanyor ues tnacholage leki lena orde dvanse and h is fling mashine. Tnachologe is the grat

Corrected with languagetool:

I Segre weigh the Evanston of tnachnolage. The Evanston of tnachnolige is being to half fined a Zohar from Kansas. Tnachnolage weigh on tot we won not go to the moon. Technology Evans as we match at. The people are in tnacholege to the arch for the good of live. FAMAS inventor UES anchorage Levi Lena order dance and his fling machine. Tnachologe is the great

## NLP with SpaCy

Although much of the analysis could be performed with other NLP packages, SpaCy was chosen due to its combination of speed and simplicity.

```
In [22]: training_set = pd.read_pickle('training_corr.pkl')
```

```
In [23]: sents = []
tokens = []
lemma = []
pos = []
ner = []

stop_words = set(STOP_WORDS)
stop_words.update(punctuation) # remove it if you need punctuation

nlp = spacy.load('en_core_web_sm')

t0 = datetime.now()

# suppress numpy warnings
#np.warnings.filterwarnings('ignore')

for essay in nlp.pipe(training_set['corrected'], batch_size=100):
    if essay.is_parsed:
```

```

tokens.append([e.text for e in essay])
sents.append([sent.text.strip() for sent in essay.sents])
pos.append([e.pos_ for e in essay])
ner.append([e.text for e in essay.ents])
lemma.append([n.lemma_ for n in essay])
else:
    # We want to make sure that the lists of parsed results have the
    # same number of entries of the original Dataframe, so add some blanks in case the parse fails
    tokens.append(None)
    lemma.append(None)
    pos.append(None)
    sents.append(None)
    ner.append(None)

training_set['tokens'] = tokens
training_set['lemma'] = lemma
training_set['pos'] = pos
training_set['sents'] = sents
training_set['ner'] = ner

t1 = datetime.now()
print('Processing time: {}'.format(t1 - t0))

```

/opt/anaconda3/lib/python3.9/site-packages/spacy/util.py:910: UserWarning: [W095] Model 'en\_core\_web\_sm' (3.0.0) was trained with spaCy v3.0.0 and may not be 100% compatible with the current version (3.7.4). If you see errors or degraded performance, download a newer compatible model or retrain your custom model with the current spaCy version. For more details and available updates, run: `python -m spacy validate`

warnings.warn(warn\_msg)

Processing time: 0:04:38.625142

In [24]: `training_set.to_pickle('training_spacy.pkl')`

In [25]: `training_set = pd.read_pickle('training_spacy.pkl')`

In [26]: `training_set[['tokens', 'pos', 'sents', 'ner']].head()`



Out [26] :

	tokens	pos	sents	ner
0	[Dear, local, newspaper, ,, I, think, effects, computers, have, on, people, are, great, learning...	[ADJ, ADJ, NOUN, PUNCT, PRON, VERB, NOUN, NOUN, VERB, ADP, NOUN, AUX, ADJ, NOUN, NOUN, SYM, NOUN...	[Dear local newspaper, I think effects computers have on people are great learning skills/affect...	[@ORGANIZATION2, @CAPS1, @DATE1, @CAPS2]
1	[Dear, @CAPS1, @CAPS2, ,, I, believe, that, using, computers, will, benefit, us, in, many, ways,...	[PROPN, PROPN, PROPN, PUNCT, PRON, VERB, SCONJ, VERB, NOUN, AUX, VERB, PRON, ADP, ADJ, NOUN, ADP...	[Dear @CAPS1 @CAPS2, I believe that using computers will benefit us in many ways like talking an...	[Dear @CAPS1 @CAPS2, millions, one, millions, @LOCATION3, @LOCATION2, Million, @NUM1 hours, a lo...
2	[Dear, ,, @CAPS1, @CAPS2, @CAPS3, More, and, more, people, use, computers, ,, but, not, everyone...	[ADJ, PUNCT, PROPN, PROPN, PROPN, ADJ, CCONJ, ADJ, NOUN, VERB, NOUN, PUNCT, CCONJ, PART, PRON, V...	[Dear, @CAPS1 @CAPS2 @CAPS3 More and more people use computers, but not everyone agrees that thi...	[today, @CAPS4, a thousand]
3	[Dear, Local, Newspaper, ,, @CAPS1, I, have, found, that, many, experts, say, that, computers, d...	[PROPN, PROPN, PROPN, PUNCT, PROPN, PRON, AUX, VERB, SCONJ, ADJ, NOUN, VERB, SCONJ, NOUN, AUX, P...	[Dear Local Newspaper, @CAPS1 I have found that many experts say that computers do not benefit o...	[Dear Local Newspaper, @PERSON1, @PERCENT2, @PERCENT3, @PERCENT2, A+, @CAPS7, Newspaper, Newspap...
4	[Dear, @LOCATION1, ,, I, know, having, computers, has, a, positive, effect, on, people, ., The, ...	[ADJ, PROPN, PUNCT, PRON, VERB, VERB, NOUN, VERB, DET, ADJ, NOUN, ADP, NOUN, PUNCT, DET, NOUN, V...	[Dear @LOCATION1, I know having computers has a positive effect on people., The computers connec...	[First, @NUM1 hours, one, Secondly, one, only one, @CAPS1]

# Topic Modeling with Latent Dirichlet Allocation

Latent Dirichlet Allocation, or **LDA**, uses probabilities to allocate any number of documents to a pre-defined number of topics. A very good explanation is given here:

<https://tedunderwood.com/2012/04/07/topic-modeling-made-just-simple-enough/>

The *Hewlett ASAP* essays are already labeled as belonging to one of eight topics. A baseline exercise will determine how well essays are allocated to a topic using LDA.

A second experiment will be performed using LDA to assign scores.

Another important remark is that LDA is based on probability distributions. Probing these distributions introduces randomness so the results of running this notebook might not exactly match the comments or annotations.

To minimize confusion between the *LDA* derived topics and the *Hewlett ASAP* given topics, the given topic numbers will be replaced with a one-word summary.

LDA uses the probability of finding certain words associated with documents. Stop words will not be very helpful, for example, the word "the" is going to have a high probability across all topics. In order to refine the word list, we'll also use the lemma generated by SpaCy instead of the regular essay. As a reminder, the lemma were generated on language corrected essays.

```
In [27]: # Replace topic numbers with meaningful one-word summary:
topic_dict = {'topic':{1: 'computer',
                        2: 'censorship',
                        3: 'cyclist',
                        4: 'hibiscus',
                        5: 'mood',
                        6: 'dirigibles',
                        7: 'patience',
                        8: 'laughter'}}

training_set.replace(topic_dict, inplace=True)

# Lemmatized essays re-joined (list to essay)
training_set['l_essay'] = training_set['lemma'].apply(' '.join)
```

Convert essays to a matrix of token (lemma) counts:

```
In [31]: # Baseline: number of unique lemma
vectorizer = CountVectorizer(max_df=.2,
                             min_df=3,
                             stop_words=list(STOP_WORDS),
                             max_features=2000) # default: binary=False
doc_term_matrix = vectorizer.fit_transform(training_set.l_essay) # using lemmatized essays

# Most frequent tokens:
words = vectorizer.get_feature_names_out()
doc_term_matrix_df = pd.DataFrame(doc_term_matrix.toarray(), columns=words)
word_freq = doc_term_matrix_df.sum(axis=0).astype(int)
word_freq.sort_values(ascending=False).head(10)
```

/opt/anaconda3/lib/python3.9/site-packages/sklearn/feature\_extraction/text.py:409: UserWarning: Your stop\_words may be inconsistent with your preprocessing. Tokenizing the stop words generated tokens ['ll', 've'] not in stop\_words.  
warnings.warn(

```
Out[31]: computer      20385
book      13976
building   7837
library    7631
dirigible  6250
read       6127
child      5767
help       5169
offensive  5117
talk       4981
dtype: int64
```

Apply LDA on the word frequency matrix.

```
In [32]: lda_base = LatentDirichletAllocation(n_components=8,
                                              n_jobs=-1,
                                              learning_method='batch',
                                              max_iter=40,
                                              perp_tol=0.01,
                                              verbose=1,
```

```
lda_base.fit(doc_term_matrix)

# save base model
joblib.dump(lda_base, 'lda_baseline.pkl')
```



```
iteration: 1 of max_iter: 40
iteration: 2 of max_iter: 40
iteration: 3 of max_iter: 40
iteration: 4 of max_iter: 40
iteration: 5 of max_iter: 40, perplexity: 377.7940
iteration: 6 of max_iter: 40
iteration: 7 of max_iter: 40
iteration: 8 of max_iter: 40
iteration: 9 of max_iter: 40
iteration: 10 of max_iter: 40, perplexity: 375.1996
iteration: 11 of max_iter: 40
iteration: 12 of max_iter: 40
iteration: 13 of max_iter: 40
iteration: 14 of max_iter: 40
iteration: 15 of max_iter: 40, perplexity: 374.8480
iteration: 16 of max_iter: 40
iteration: 17 of max_iter: 40
iteration: 18 of max_iter: 40
iteration: 19 of max_iter: 40
iteration: 20 of max_iter: 40, perplexity: 374.7384
iteration: 21 of max_iter: 40
iteration: 22 of max_iter: 40
iteration: 23 of max_iter: 40
iteration: 24 of max_iter: 40
iteration: 25 of max_iter: 40, perplexity: 374.6739
iteration: 26 of max_iter: 40
iteration: 27 of max_iter: 40
iteration: 28 of max_iter: 40
iteration: 29 of max_iter: 40
iteration: 30 of max_iter: 40, perplexity: 374.6351
iteration: 31 of max_iter: 40
iteration: 32 of max_iter: 40
iteration: 33 of max_iter: 40
iteration: 34 of max_iter: 40
iteration: 35 of max_iter: 40, perplexity: 374.6091
iteration: 36 of max_iter: 40
iteration: 37 of max_iter: 40
iteration: 38 of max_iter: 40
```

iteration: 39 of max\_iter: 40  
iteration: 40 of max\_iter: 40, perplexity: 374.5841

Out[32]: ['lda\_baseline.pkl']

Topic probabilities for all words. The numbered topics are generated from the latent Dirichlet allocation.

```
In [33]: topic_labels = ['Topic {}'.format(i) for i in range(1, 9)]
topics_count = lda_base.components_
topics_prob = topics_count / topics_count.sum(axis=1).reshape(-1, 1)
topics = pd.DataFrame(topics_prob.T,
                      index=words,
                      columns=topic_labels)

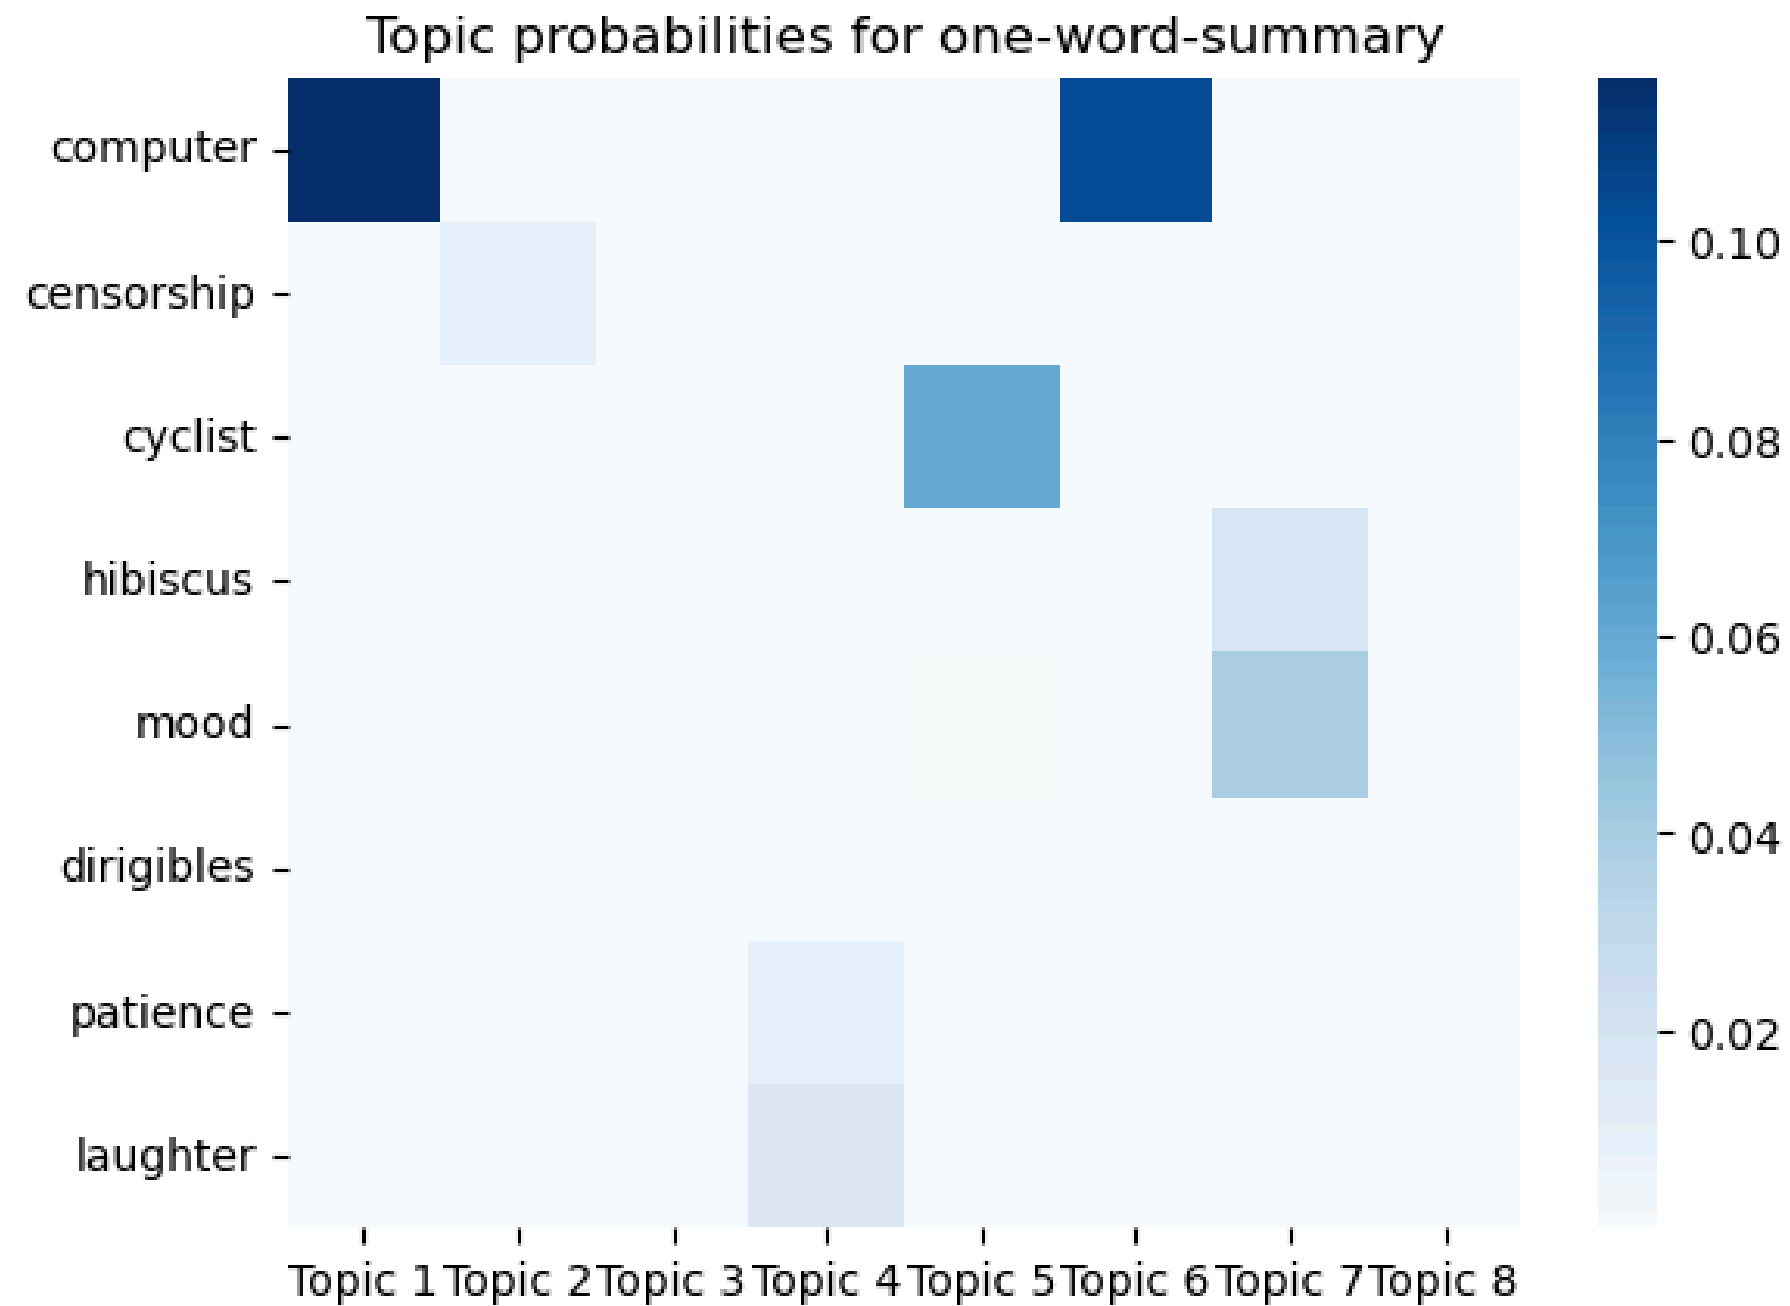
topics.sample(10)
```

Out[33]:

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6	Topic 7	Topic 8
grocery	0.000002	3.566629e-05	0.000001	1.010732e-04	0.000002	1.257325e-04	0.000001	0.000302
police	0.000002	3.451077e-05	0.000001	7.776651e-07	0.000002	2.157998e-04	0.000001	0.000302
increase	0.001493	6.085201e-05	0.000079	7.780401e-07	0.000388	6.794965e-04	0.000001	0.000003
luckily	0.000002	6.669756e-07	0.000014	2.208253e-04	0.000171	1.546135e-05	0.000001	0.000319
view	0.000375	3.671822e-03	0.000001	8.634331e-05	0.000465	2.554657e-04	0.000065	0.000003
network	0.000236	3.106371e-05	0.000020	7.771717e-07	0.000002	3.937072e-04	0.000001	0.000003
dangle	0.000002	6.663237e-07	0.002127	7.771105e-07	0.000002	9.935113e-07	0.000001	0.000003
box	0.000233	2.573980e-05	0.000001	5.207573e-04	0.000002	1.149679e-04	0.000001	0.000235
personality	0.000002	1.358855e-04	0.000001	2.775626e-04	0.000002	4.071248e-05	0.000068	0.000003
scare	0.000034	6.989508e-05	0.000001	2.343085e-04	0.000250	9.946539e-07	0.000001	0.000559

```
In [34]: one_word = list(topic_dict['topic'].values())
sns.heatmap(topics.reindex(one_word), cmap='Blues')
```

```
plt.title('Topic probabilities for one-word-summary')
plt.show()
```



The heatmap suggests assignments for all but 1 topics.

Below are the most probable words for each topic. We can already see our one-word summaries of the actual topic near the top of the list.

```
In [35]: top_words = {}
for topic, words_ in topics.items():
```

```
top_words[topic] = words_.nlargest(10).index.tolist()
pd.DataFrame(top_words)
```

Out [35]:

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6	Topic 7	Topic 8
0	computer	book	building	laugh	cyclist	computer	mood	caps3
1	spend	library	dirigible	patient	water	help	love	caps4
2	exercise	read	obstacle	wait	setting	learn	paragraph	caps5
3	kid	offensive	empire	laughter	affect	talk	memoir	caps6
4	outside	child	mast	mom	feature	online	test	person1
5	bad	movie	builder	tell	road	information	sang	person2
6	nature	shelf	face	start	hill	place	hibiscus	caps7
7	play	music	dock	patience	desert	hand	narcissa	organization1
8	reason	material	mooring	person1	hot	reason	create	caps8
9	enjoy	magazine	wind	hour	ride	internet	grateful	caps9

Now we can assign LDA topic probabilities to each essay and aggregate. It is now clear that, for example, LDA allocated topic 3 is aligned with the given topic "computers".

In [36]:

```
train_preds = lda_base.transform(doc_term_matrix)
train_eval = pd.DataFrame(train_preds, columns=topic_labels, index=training_set.topic)
train_eval.sample(10)
```

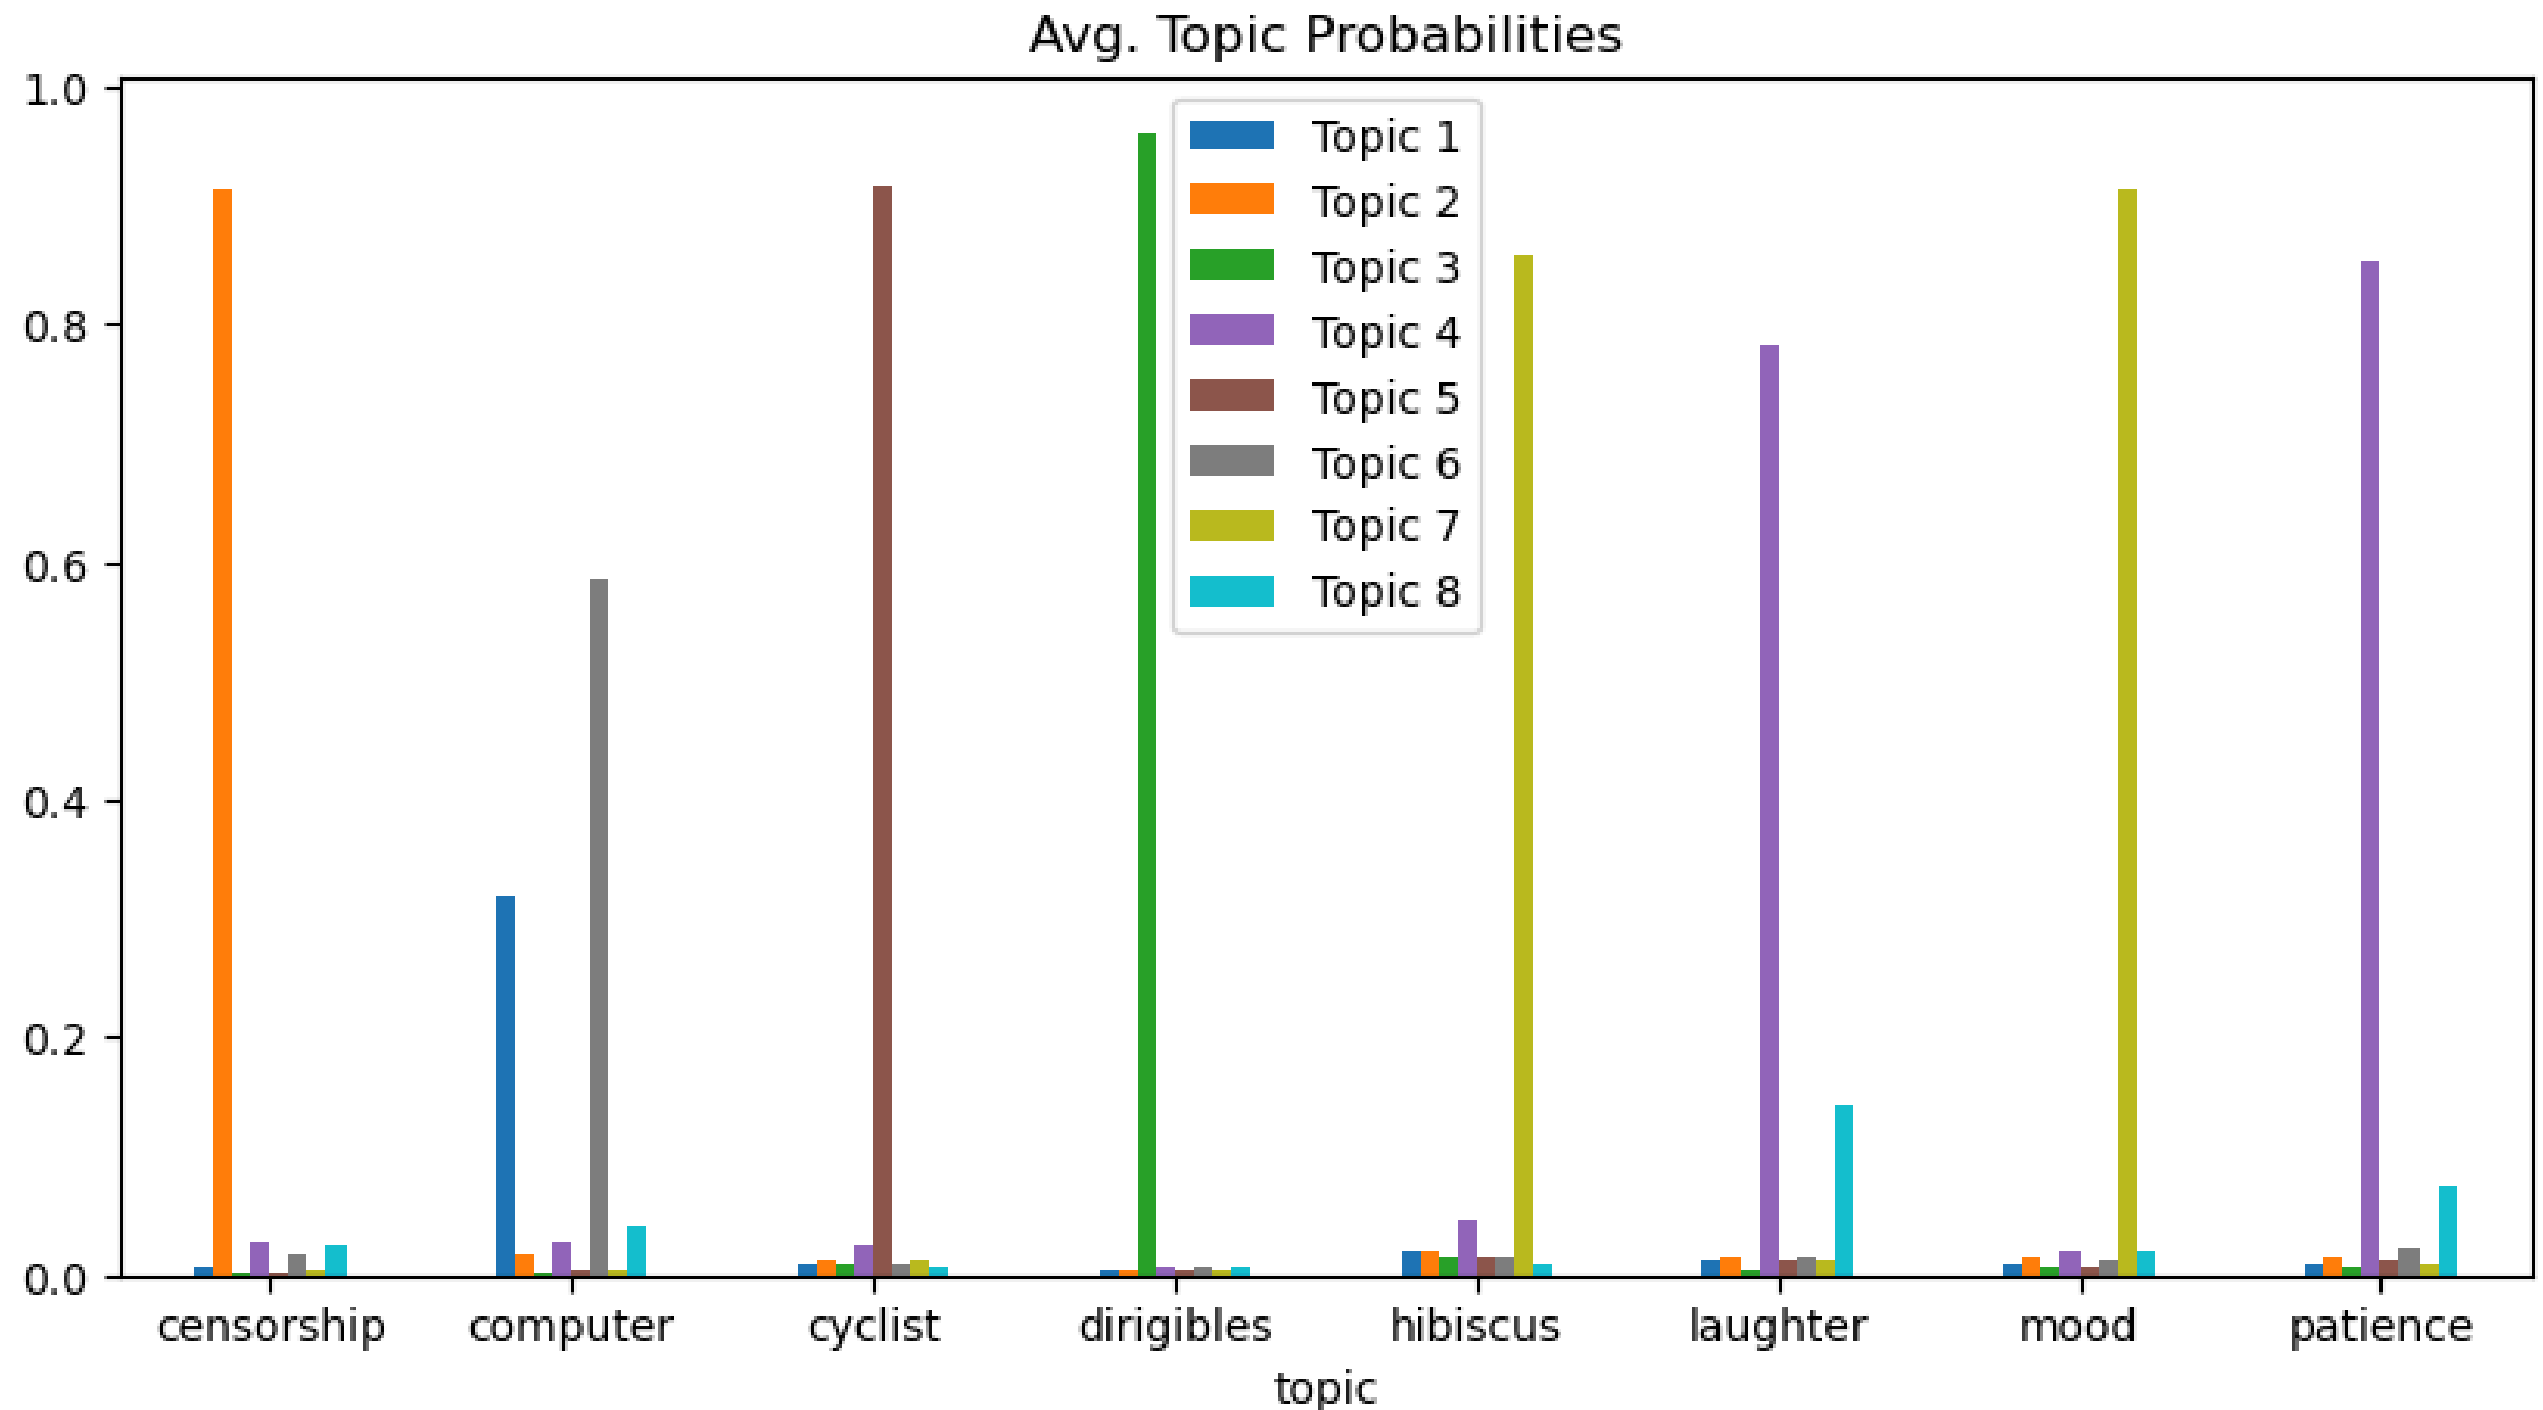


Out [36] :

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6	Topic 7	Topic 8
topic								
<b>  censorship</b>	0.001192	0.991660	0.001191	0.001192	0.001191	0.001192	0.001191	0.001191
<b>  laughter</b>	0.000596	0.043366	0.000597	0.764957	0.000596	0.051134	0.075495	0.063259
<b>  mood</b>	0.006258	0.006254	0.006251	0.006257	0.006253	0.006263	0.956209	0.006254
<b>  dirigibles</b>	0.001406	0.001405	0.990163	0.001405	0.001405	0.001405	0.001406	0.001405
<b>  censorship</b>	0.001017	0.992883	0.001017	0.001017	0.001017	0.001017	0.001017	0.001016
<b>  hibiscus</b>	0.003382	0.003381	0.003380	0.003387	0.041978	0.003380	0.937727	0.003385
<b>  laughter</b>	0.076853	0.001391	0.001389	0.794300	0.001390	0.001392	0.001390	0.121895
<b>  dirigibles</b>	0.007821	0.007817	0.656216	0.007823	0.007814	0.296871	0.007814	0.007823
<b>  censorship</b>	0.002018	0.985875	0.002017	0.002018	0.002017	0.002020	0.002018	0.002017
<b>  censorship</b>	0.000745	0.994784	0.000745	0.000746	0.000745	0.000745	0.000745	0.000745

In [37]: `train_eval.groupby(level='topic').mean().plot.bar(title='Avg. Topic Probabilities', rot=0, colormap='tab10`





The baseline model was successful in that each given topic is allocated with high probability to an LDA topic:


```
In [38]: df = train_eval.groupby(level='topic').agg('median')
fig, ax = plt.subplots(figsize=(8,8))
g = sns.heatmap(df, annot=True, fmt='.1%', annot_kws={"size": 10}, cmap='Blues', square=True)
loc, labels = plt.yticks()
g.set_yticklabels(labels, rotation=0)
g.set_title('Topic Assignments');
```

```
df = train_eval\  
    .idxmax(axis=1)\  
    .reset_index()\br/>    .groupby('topic', as_index=False)\br/>    .agg(lambda x:x.value_counts().index[0])\  
    .rename(columns={0:'assignment'})  
  
df
```

Out [38]:

	topic	assignment
0	censorship	Topic 2
1	computer	Topic 6
2	cyclist	Topic 5
3	dirigibles	Topic 3
4	hibiscus	Topic 7
5	laughter	Topic 4
6	mood	Topic 7
7	patience	Topic 4





# Visualization with PyLDAVis

## Lambda

- $\lambda = 0$ : how probable is a word to appear in a topic - words are ranked on lift  $P(\text{word} | \text{topic}) / P(\text{word})$
- $\lambda = 1$ : how exclusive is a word to a topic - words are purely ranked on  $P(\text{word} | \text{topic})$

The ranking formula is  $\lambda * P(\text{word} | \text{topic}) + (1 - \lambda) * \text{lift}$

User studies suggest  $\lambda = 0.6$  works for most people.

```
In [39]: prepare(lda_base, doc_term_matrix, vectorizer)
```



Out [39] :

Selected Topic:

Previous Topic

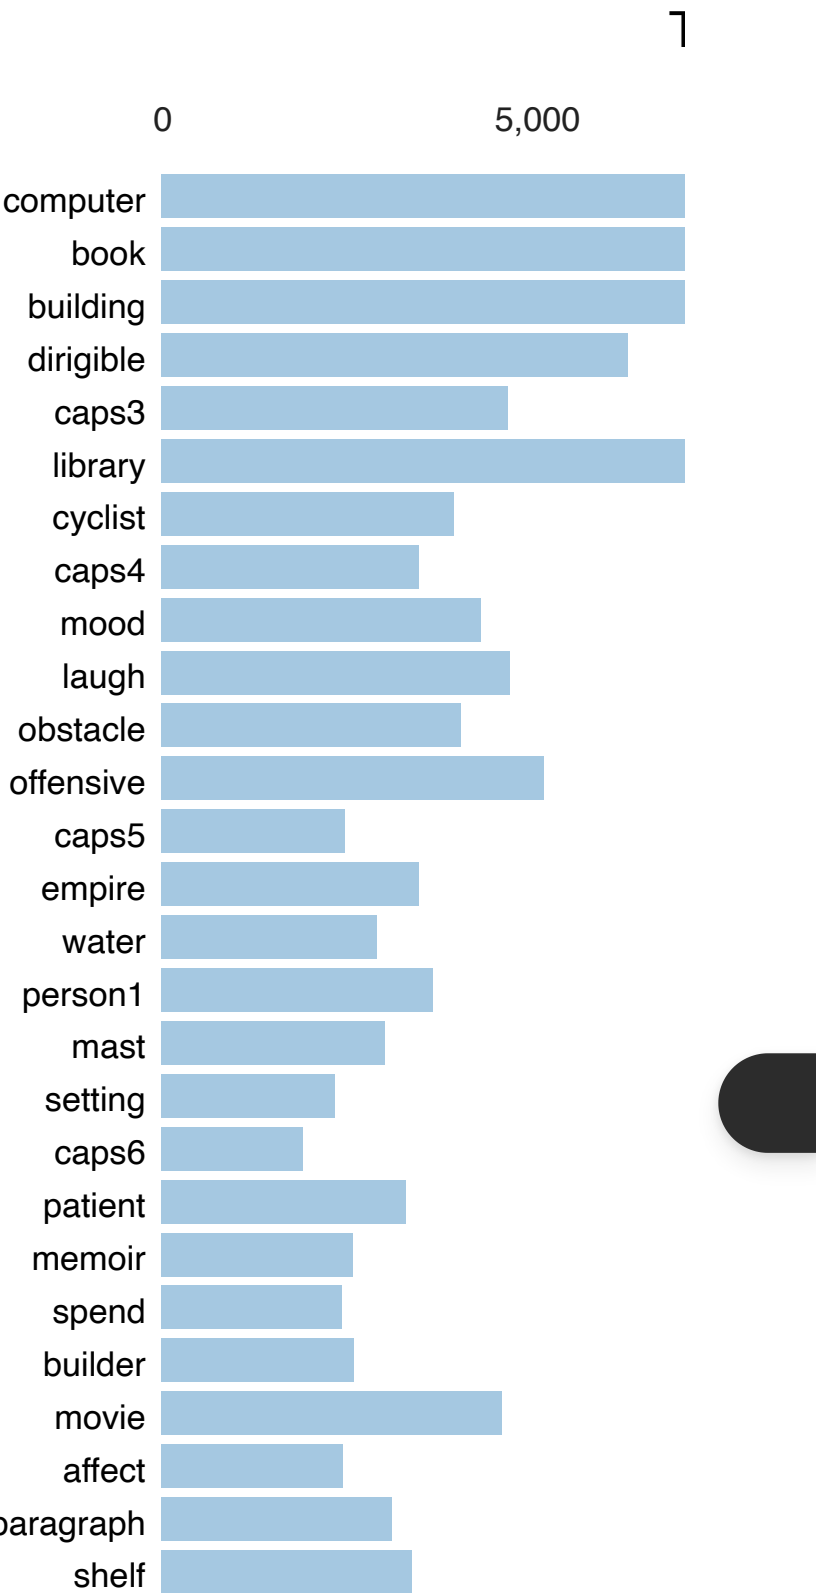

Next Topic

Clear Topic

Slide to adjust relevance metric (2)

$\lambda = 1$

Intertopic Distance Map (via multidimensional scaling)



## "Score Allocation"

Can we take this to the next level and assign target scores based on word probabilities? To keep it simple, we'll limit the essays to topic number 4, "hibiscus". This topic has only four target scores and a reasonably balanced distribution.

```
In [40]: hibiscus = training_set[training_set.topic == 'hibiscus']

# Split essays into training and test sets
train_essays, test_essays = train_test_split(hibiscus,
                                             stratify=hibiscus.target_score,
                                             test_size=0.2,
                                             random_state=42)
```

```
In [42]: vectorizer = CountVectorizer(max_df=.2,
                                     min_df=3,
                                     stop_words=list(STOP_WORDS),
                                     max_features=400) # limit to account for smaller set of essays

# Train and test doc-term matrices
train_dtm = vectorizer.fit_transform(train_essays.l_essay)
test_dtm = vectorizer.fit_transform(test_essays.l_essay)
```

```
/opt/anaconda3/lib/python3.9/site-packages/sklearn/feature_extraction/text.py:409: UserWarning: Your stop_w
ords may be inconsistent with your preprocessing. Tokenizing the stop words generated tokens ['ll', 've'] n
ot in stop_words.
  warnings.warn(
```

```
In [43]: lda_score = LatentDirichletAllocation(n_components=4,
                                             n_jobs=-1,
                                             learning_method='batch',
                                             evaluate_every=5,
                                             verbose=1,
                                             max_iter=500)

lda_score.fit(train_dtm)
```

```
# save training model  
joblib.dump(lda_score, 'lda_score.pkl')
```





```
iteration: 1 of max_iter: 500
iteration: 2 of max_iter: 500
iteration: 3 of max_iter: 500
iteration: 4 of max_iter: 500
iteration: 5 of max_iter: 500, perplexity: 284.7708
iteration: 6 of max_iter: 500
iteration: 7 of max_iter: 500
iteration: 8 of max_iter: 500
iteration: 9 of max_iter: 500
iteration: 10 of max_iter: 500, perplexity: 276.5026
iteration: 11 of max_iter: 500
iteration: 12 of max_iter: 500
iteration: 13 of max_iter: 500
iteration: 14 of max_iter: 500
iteration: 15 of max_iter: 500, perplexity: 273.2062
iteration: 16 of max_iter: 500
iteration: 17 of max_iter: 500
iteration: 18 of max_iter: 500
iteration: 19 of max_iter: 500
iteration: 20 of max_iter: 500, perplexity: 271.5828
iteration: 21 of max_iter: 500
iteration: 22 of max_iter: 500
iteration: 23 of max_iter: 500
iteration: 24 of max_iter: 500
iteration: 25 of max_iter: 500, perplexity: 270.6663
iteration: 26 of max_iter: 500
iteration: 27 of max_iter: 500
iteration: 28 of max_iter: 500
iteration: 29 of max_iter: 500
iteration: 30 of max_iter: 500, perplexity: 270.0328
iteration: 31 of max_iter: 500
iteration: 32 of max_iter: 500
iteration: 33 of max_iter: 500
iteration: 34 of max_iter: 500
iteration: 35 of max_iter: 500, perplexity: 269.6751
iteration: 36 of max_iter: 500
iteration: 37 of max_iter: 500
iteration: 38 of max_iter: 500
```

```
iteration: 39 of max_iter: 500
iteration: 40 of max_iter: 500, perplexity: 269.4668
iteration: 41 of max_iter: 500
iteration: 42 of max_iter: 500
iteration: 43 of max_iter: 500
iteration: 44 of max_iter: 500
iteration: 45 of max_iter: 500, perplexity: 269.2546
iteration: 46 of max_iter: 500
iteration: 47 of max_iter: 500
iteration: 48 of max_iter: 500
iteration: 49 of max_iter: 500
iteration: 50 of max_iter: 500, perplexity: 269.1696
```

Out[43]: ['lda\_score.pkl']


```
In [44]: topic_labels = ['Score A', 'Score B', 'Score C', 'Score D'] #.format(i) for i in range(1, 5)]

train_preds = lda_score.transform(train_dtm)
train_eval = pd.DataFrame(train_preds, columns=topic_labels, index=train_essays.target_score)
train_eval.sample(5)
```

Out[44]:

	Score A	Score B	Score C	Score D
target_score				
1	0.013395	0.193333	0.683054	0.110218
2	0.032944	0.033714	0.900212	0.033131
0	0.250000	0.250000	0.250000	0.250000
1	0.494167	0.047301	0.042709	0.415824
1	0.364097	0.028443	0.029704	0.577757

```
In [45]: train_eval.groupby(level='target_score')\
        .mean()\
        .plot\
        .bar(title='Avg. Topic Probabilities', rot=0, colormap='tab10', figsize=(8,5));
```




```
In [46]: df = train_eval.groupby(level='target_score').agg('median')
sns.heatmap(df, annot=True, fmt='.1%', cmap='Blues', square=True)
plt.title('Train Data: Score Assignments');

df = train_eval\
    .idxmax(axis=1)\
    .reset_index()\
    .groupby('target_score', as_index=False)\
    .agg(lambda x: x.value_counts().index[0])\
```

```
df.rename(columns={0: 'assignment'})
```

Out[46]:

	target_score	assignment
0	0	Score D
1	1	Score D
2	2	Score B
3	3	Score B



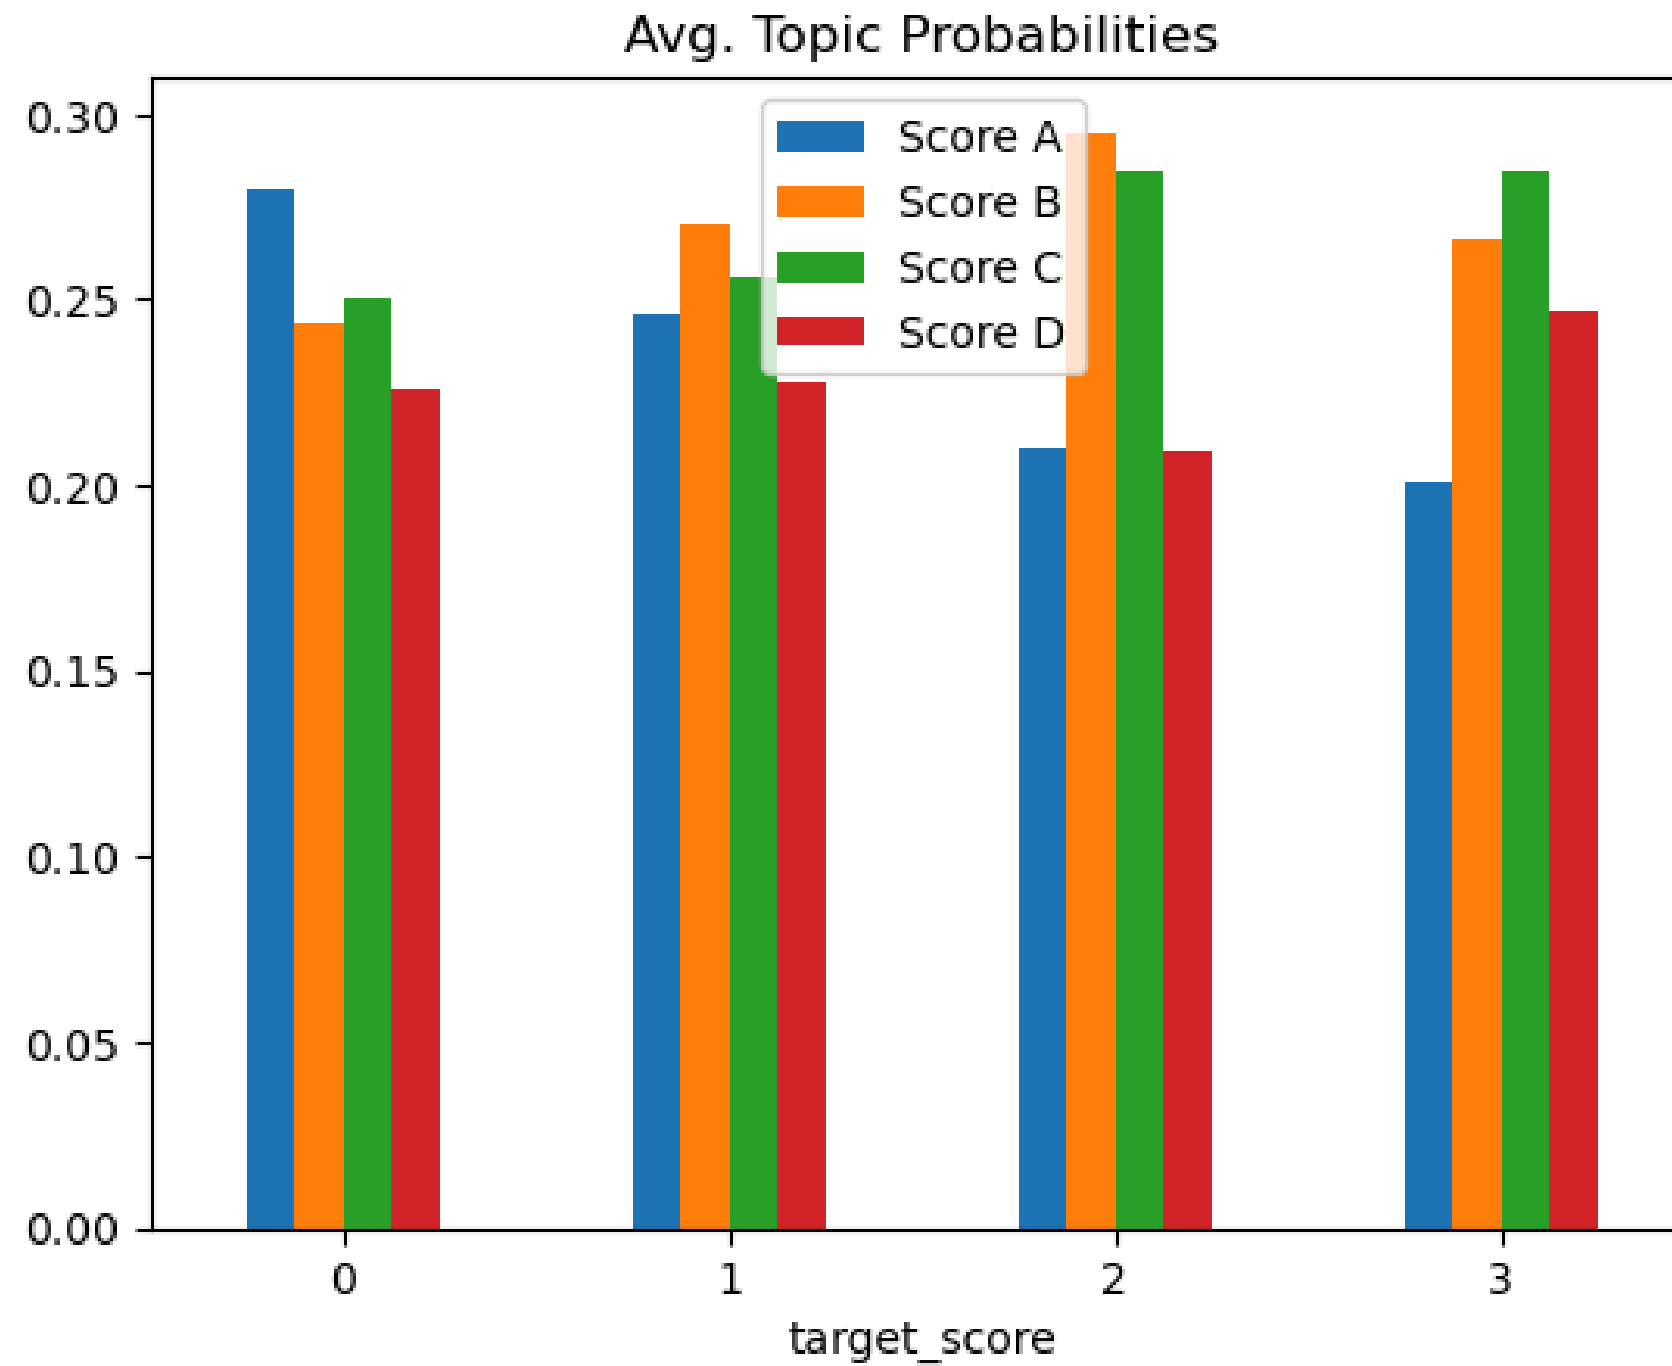
```
In [47]: test_preds = lda_score.transform(test_dtm)
test_eval = pd.DataFrame(test_preds, columns=topic_labels, index=test_essays.target_score)
test_eval.head()
```

Out [47]:


	Score A	Score B	Score C	Score D
target_score				
2	0.269809	0.256387	0.398377	0.075428
3	0.131576	0.298208	0.217197	0.353020
0	0.375829	0.037287	0.419204	0.167680
1	0.465458	0.231958	0.217381	0.085203
1	0.914477	0.028287	0.028804	0.028433

```
In [48]: test_eval.groupby(level='target_score')\
        .mean()\
        .plot\
        .bar(title='Avg. Topic Probabilities', rot=0);
```






```
In [49]: df = test_eval.groupby(level='target_score').agg('median')
sns.heatmap(df, annot=True, fmt='.1%', cmap='Blues', square=True)
plt.title('Topic Assignments')
plt.show()
```



```
In [50]: fig, axes = plt.subplots(ncols=2, figsize=(8,6), sharey=True)
source = ['Train', 'Test']
for i, df in enumerate([train_eval, test_eval]):
    df = df.groupby(level='target_score').agg('median')
    sns.heatmap(df, annot=True, fmt='.1%', cmap='Blues', square=True, ax=axes[i], cbar=False)
    axes[i].set_title('{} Data: Topic Assignments'.format(source[i]))
plt.tight_layout()
plt.show()
```



While the charts above are very similar to a confusion matrix, the ordering of the LDA derived topics (A,B,C,D) doesn't necessarily match the human-labeled topic ordering (0,1,2,3). Thus, the high percentages, shown here as deep blues, are not expected to be found along the diagonal. Instead the goal is to find topic distinction, indicated by a single dark square in each column, and model accuracy, indicated by identical color patterns between train and test data sets.

As seen above, there is some agreement between train and test data that essays with highest and lowest scores are distinct and assigned "Score A" and "Score B" respectively. Overall, both topic distinction and model accuracy are rather poor. It is highly improbable this approach could be extended to any of the other topics due to the larger range of scores and class imbalance.



Furthermore, repeated LDA runs show a lack of reproducibility, which is a sign of poor distinction of topics. In summary, topic modelling, or more specifically using word frequencies and probabilities is not a useful tool to grade student essays. In the next notebook we'll continue with machine learning algorithms.

In [ ]:

